



## Final Project Report

Two-Tier web application automation with Terraform,  
Ansible and GitHub Actions

Ammar Salmawy - 110561198  
Behzad Rajabalipour - 160934212

Prof: Dhana Karuppusamy  
ASC730

Date: 9 Dec 2024

## Introduction

In this project, we will use Terraform, Ansible, and GitHub actions to deploy and configure a highly available website. First, Terraform will be used to deploy the networking infrastructure which includes VPC, route tables, security groups, and nat/internet gateways then we will deploy and configure the EC2 instances to host the website, auto-scaling group and load balancer. Two of the EC2 instances will be configured as web servers using user data as part of the EC2 initialization. Second, the Ansible playbook will be used to configure the Apache servers for the rest of the instances. Lastly, we will use GitHub and GitHub actions as a continuous integration tool to automate the deployment of the website.

## Traffic flow Terraform

This report will share snippets of the code and explain its functionality and what goal it serves as the overall deployment will be in the demo recording.

First Networking deployment using Terraform

```
module "vpc-dev" {  
  # source = "../..//prod_network"  
  source      = "git::https://github.com/Behzad-Rajabalipour/prod_network.git" # it's be  
  env         = var.env  
  vpc_cidr    = var.vpc_cidr  
  public_cidr_blocks = var.public_cidr_blocks  
  private_cidr_blocks = var.private_cidr_blocks  
  prefix      = var.prefix  
  default_tags = var.default_tags  
}
```

*We used GitHub repo as the source of our code*

```

# Create a new VPC
resource "aws_vpc" "main" {
  cidr_block      = var.vpc_cidr
  instance_tenancy = "default"
  tags = merge(
    local.default_tags, {
      Name = "VPC-${var.prefix}"
    }
  )
}

```

Creating a new VPC

```

# Create Internet Gateway
resource "aws_internet_gateway" "igw" {
  vpc_id = aws_vpc.main.id

  tags = merge(local.default_tags,
    {
      "Name" = "IGW-${var.prefix}"
    }
  )
}

# Create the NAT Gateway
resource "aws_nat_gateway" "nat_gw" {
  allocation_id = aws_eip.nat_eip.id
  subnet_id     = aws_subnet.public_subnet[0].id # Ensure this is your public subnet
  tags = {
    Name = "${var.prefix}-nat-gateway"
  }
}

```

Create an internet gateway and NAT gateway for internet access

The NAT gateway is deployed into a public subnet and its functionality is to provide access to the internet for the EC2 instances in the Private subnet to install the Apache server.

```

# Add provisioning of the public subnet in the default VPC
resource "aws_subnet" "public_subnet" {
  count          = length(var.public_cidr_blocks)
  vpc_id        = aws_vpc.main.id
  cidr_block     = var.public_cidr_blocks[count.index]
  availability_zone = data.aws_availability_zones.available.names[count.index] # so these two subnets are in
  tags = merge(
    local.default_tags, {
      Name = "${var.prefix}-public-subnet-${count.index + 1}"
    }
  )
}

```

```

# Provision public subnets in custom VPC
variable "public_cidr_blocks" {
  default     = ["10.1.1.0/24", "10.1.2.0/24", "10.1.3.0/24", "10.1.4.0/24"]
  type       = list(string)
  description = "Public Subnet CIDRs"
}

```

Deploying 4 public subnets into 4 different availability zones

```
resource "aws_subnet" "private_subnet" {
  count          = length(var.private_cidr_blocks)
  vpc_id         = aws_vpc.main.id
  cidr_block     = var.private_cidr_blocks[count.index]
  availability_zone = data.aws_availability_zones.available.names[count.index] #
  tags = merge(
    local.default_tags, {
      Name = "${var.prefix}-private-subnet-${count.index + 1}"
    }
  )
}
```

```
31
32 # Provision private subnets in custom VPC
33 variable "private_cidr_blocks" {
34   default     = ["10.1.5.0/24", "10.1.6.0/24"]
35   type        = list(string)
36   description = "Public Subnet CIDRs"
37 }
38
```

Deploying 2 private subnets into the first 2 availability zones

```
38 # Route table to route add default gateway pointing to Internet Gateway (IGW)
39 resource "aws_route_table" "public_route_table" {
40   vpc_id = aws_vpc.main.id
41   route {
42     cidr_block = "0.0.0.0/0"
43     gateway_id = aws_internet_gateway.igw.id
44   }
45   tags = {
46     Name = "${var.prefix}-RT-public-subnets"
47   }
48 }
49
50 # Route table for private subnets pointing to NAT Gateway
51 resource "aws_route_table" "private_route_table" {
52   vpc_id = aws_vpc.main.id
53   route {
54     cidr_block = "0.0.0.0/0"
55     gateway_id = aws_nat_gateway.nat_gw.id
56   }
57   tags = {
58     Name = "${var.prefix}-RT-private-subnets"
59   }
60 }
```

Traffic is routed to the IGW for public subnets and to the NAT gateway from the private subnet

```

# Load Balancer
resource "aws_lb" "ALB" {
  name                = "ALB"
  internal            = false
  load_balancer_type = "application"
  security_groups     = [aws_security_group.lb_sg.id]
  subnets            = aws_subnet.public_subnet[*].id # ids of both p

  enable_deletion_protection = false

  tags = {
    Name = "MyApplicationLoadBalancer"
  }
}

```

Deploy the load balancer into two subnets

```

# Security Group for the Load Balancer
resource "aws_security_group" "lb_sg" {
  name        = "lb_sg"
  vpc_id      = aws_vpc.main.id
  description = "Allow HTTP and HTTPS traffic"

  ingress {
    from_port = 80
    to_port   = 80
    protocol  = "tcp"
    cidr_blocks = ["0.0.0.0/0"] # Allow HTTP from anywhere
  }

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1" # Allow all outbound traffic
    cidr_blocks = ["0.0.0.0/0"]
  }
}

```

Allow inbound traffic to port 80(HTTP) for the load balancer

```

# Listener for the Load Balancer
resource "aws_lb_listener" "listener" {
  load_balancer_arn = aws_lb.ALB.arn
  port              = 80
  protocol           = "HTTP"

  default_action {
    type = "forward"
    target_group_arn = aws_lb_target_group.TG.arn
  }
}

```

Create an HTTP listener on the ALB that forwards traffic to the target group

```

1 # it only can get output.tf of networkmodule
2 output "public_subnet_ids" {
3   value = module.vpc-dev.public_subnet_ids # error => value = aws_subnet.public_subnet[*].id
4 }
5
6 # it only can get output.tf of networkmodule
7 output "public_subnet_cidrs" {
8   value = module.vpc-dev.public_subnet_cidrs # error => value = aws_subnet.public_subnet[*].id
9 }
10
11 # it only can get output.tf of networkmodule
12 output "private_subnet_ids" {
13   value = module.vpc-dev.private_subnet_ids # error => value = aws_subnet.private_subnet[*].id
14 }
15
16 output "public_route_table_id" {
17   value = module.vpc-dev.public_route_table_id
18 }
19
20 output "vpc_id" {
21   value = module.vpc-dev.vpc_id # error => value = aws_vpc.main.id
22 }
23
24 output "vpc_cidr" {
25   value = module.vpc-dev.vpc_cidr
26 }
27
28 output "target_group_arn" {
29   value = module.vpc-dev.target_group_arn
30 }
31
32 output "ALB_SG_id" {
33   value = module.vpc-dev.ALB_SG_id
34 }
35
36

```

Output subnets, ABL, TG and VPC IDs so it can be used when deploying the web servers

```

47 resource "aws_instance" "public_instance" {
48   count                = length(data.terraform_remote_state.prod_net_tfstate.outputs.public_subnet_ids)
49   ami                 = data.aws_ami.latest_amazon_linux.id
50   instance_type       = lookup(var.instance_type, var.env)
51   key_name            = aws_key_pair.keyName.key_name # because we used count in key_name so we have to give it's index
52   security_groups     = [aws_security_group.public_instance_sg.id]
53   subnet_id           = data.terraform_remote_state.prod_net_tfstate.outputs.public_subnet_ids[count.index] # it use s3/dev/network/terraform.tfstate => outputs.
54   associate_public_ip_address = true
55   user_data            = count.index == 0 || count.index == 2 ? templatefile("${path.module}/install_httpd.sh.tpl",
56     {
57       env      = upper(var.env),
58       prefix   = upper(var.prefix),
59       name     = "Behzad Rajabaliipour"
60     }) : null
61
62   root_block_device { # it will encrypt it if it is in envtest
63     encrypted = var.env == "test" ? true : false
64   }
65
66   lifecycle { # lifecycle will create a new instance before destroy previous one
67     create_before_destroy = true
68   }
69
70   # Dynamic tags
71   tags = merge(
72     local.default_tags,
73     count.index == 1
74     ? {
75       Name = "Bastion VM webserver${count.index + 1}",
76       Role = "Bastion VM"
77       Type = "type1"
78     }
79     : {
80       Name = "webserver${count.index + 1}",
81       Type = contains([3], count.index) ? "type1" : null
82     }
83   )
84 }
85

```

Deployment of EC2 instance into the public subnets

Using a count to iterate to all public subnets then using user data to install httpd to the first and third EC2 instances index 0 and 2. then using type1 tag for instances 2 and 4 (index 1 and 3).

The tag will be used to target these instances and install/configure the web server using Ansible

```

83
84  lifecycle { # lifecycle will create a new instance before destroy previous one
85    create_before_destroy = true
86  }
87

```

Terraform ensures a new resource is created before an existing resource is destroyed

This ensures the high availability of the websites it avoids service interruptions by maintaining the old resource until the new one is fully provisioned and ready

```

resource "aws_security_group" "private_instance_SG" {
  name           = "allow_ssh_icmp"
  description    = "Allow SSH and ICMP traffic from a specific VM"
  vpc_id         = data.terraform_remote_state.prod_net_tfstate.outputs.vpc_id

  # Ingress rule for SSH
  ingress {
    description = "Allow SSH from specific VM"
    from_port   = 22
    to_port     = 22
    protocol    = "tcp"
    cidr_blocks = ["${aws_instance.public_instance[1].private_ip}/32"] # private_ip(IPV4 pri
  }

  # Egress rule to allow all outbound traffic
  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1" # Allow all egress traffic
    cidr_blocks = ["0.0.0.0/0"]
  }

  tags = {
    Name = "Allow SSH and ICMP Security Group"
  }
}

```

Using the IP address of web server 2 to allow SSH into the EC2 instances in the private subnet

```

# Register EC2 Instances to Target Group
resource "aws_lb_target_group_attachment" "tg_attachment" {
  count          = length(aws_instance.public_instance[*].id) - 1 # exclude the last public instance (We
  target_group_arn = data.terraform_remote_state.prod_net_tfstate.outputs.target_group_arn
  target_id       = aws_instance.public_instance[count.index].id
  port           = 80
}

```

Attach the first 3 web servers to the target group

```

resource "aws_autoscaling_policy" "cpu_scale_out" {
  name                = "target-tracking-scale-up-policy"
  autoscaling_group_name = aws_autoscaling_group.public_instance_asg.name
  policy_type          = "TargetTrackingScaling"

  target_tracking_configuration {
    target_value = 60 # Trigger scale-out when CPU usage exceeds 60%
    predefined_metric_specification {
      predefined_metric_type = "ASGAverageCPUUtilization"
    }
    disable_scale_in = true # Disable scaling in
  }
}

#-----
# Simple Scaling policy
# we can have many simple scaling policy

# create cloudwatch alarm with CPUUtilization metric
resource "aws_cloudwatch_metric_alarm" "scale_in_alarm" {
  alarm_name        = "$scale-in-alarm"
  comparison_operator = "LessThanOrEqualToThreshold" # <=

  evaluation_periods = 2 # data points, Number of periods to evaluate
  datapoints_to_alarm = 2 # data points, Number of data points that must breach the threshold

  metric_name = "CPUUtilization" # metric name for this alarm
  namespace = "AWS/EC2"
  period = 60 # 60 sec for alarm interval
  statistic = "Average"
  threshold = 40 # CPUUtilization < 40%

  alarm_description = "Trigger scale-in when CPU utilization is less than 40%"
  dimensions = {
    AutoScalingGroupName = aws_autoscaling_group.public_instance_asg.name
  }
  actions_enabled = true

  alarm_actions = [aws_autoscaling_policy.scale_in_policy.arn] # Link to scaling policy
}

# create simple scaling policy
resource "aws_autoscaling_policy" "scale_in_policy" {
  name                = "$simple-scaling-in-policy"
  scaling_adjustment = -1 # Reduce by 1 instance
  adjustment_type     = "ChangeInCapacity" # Adjust the number of instances
  cooldown            = 70 # Wait 70 seconds between scaling activities
  autoscaling_group_name = aws_autoscaling_group.public_instance_asg.name
}

```

Creating a scale-in and scale-out policy for the auto-scaling group  
 This also creates an alarm for the cloud watch to monitor the EC2 instance

## Ansible traffic flow

Ansible will be used to configure instances 2 and 4 (tagged: type1) we will use playbook and jinja2 templates for this purpose.



```

plugin: aws_ec2
regions:
  - us-east-1
keyed_groups:
  - key: tags.Type
    prefix: tag          # Adds the prefix tag_ to the group names, like tag_
filters:
  instance-state-name : running
compose:
  ansible_host: public_ip_address    # ansible_host: is mandatory like ansi
  ansible_hostname: public_ip_address # i used it in index.j2 file
  # ansible_user          : ec2-user
  # ansible_ssh_private_key_file : ~/.ssh/prodKey    # it will store in the

```

Dynamic Inventory Setup: this will group instances based on their tags

```

- hosts: tag_type1
  gather_facts: True
  become: yes

  vars:
    source_file: ./index.j2
    dest_file: /var/www/html

```

Targeting type in the playbook

The screenshot displays the AWS Management Console for the 'us-east-1' region. The 'Instances' page shows a list of 7 EC2 instances. The first instance, 'Bastion VM webserver2', is selected, and its details are shown in the lower pane. The details pane shows the instance's configuration, including its name, environment, type, application, role, and owner.

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS
Bastion VM webserver2	i-002b386a0b3d5aef1	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	-
VM6	i-04928281058538f6a	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	-
webserver3	i-0e0146c6068ca05e9	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1c	-
prod-asg-instance	i-063677b55a781c136	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1a	-
webserver1	i-0bccf5d411820889b	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1a	-
Webserver5	i-0caa23216a2b0a114	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1a	-
webserver4	i-01b36bec4f08fbb4	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1d	-

Details for i-002b386a0b3d5aef1 (Bastion VM webserver2):

Key	Value
Name	Bastion VM webserver2
env	prod
Type	type1
App	FinalProject
Role	Bastion VM
Owner	Behzad

As you see in the picture, Webserver2 and 4 has a tag name: Type = type1  
So ansible inventory will dynamically take these two instances IP

```
tasks:
- name: Install Apache Web Server for RPM
  yum: name=httpd state=latest
  when: ansible_os_family == "RedHat"

- name: Print Linux Family
  debug: var=ansible_os_family

- name: Get the local IPv4 address
  shell: curl http://169.254.169.254/latest/meta-data/local-ipv4
  register: myip
  args:
    warn: false
```

Ansible playbook tasks to install httpd

```

- name: Generate index.html from jinja2 template and copy to the remote host
  template: src={{ source_file }} dest={{ dest_file }}/index.html mode=0555
  notify: Restart Httpd
  when: ansible_os_family == "RedHat"

- name: Start Apache Web Server
  service: name=httpd state=started enabled=yes
  when: ansible_os_family == "RedHat"

handlers:
- name: Restart Httpd
  service: name=httpd state=restarted
  when: ansible_os_family == "RedHat"

```

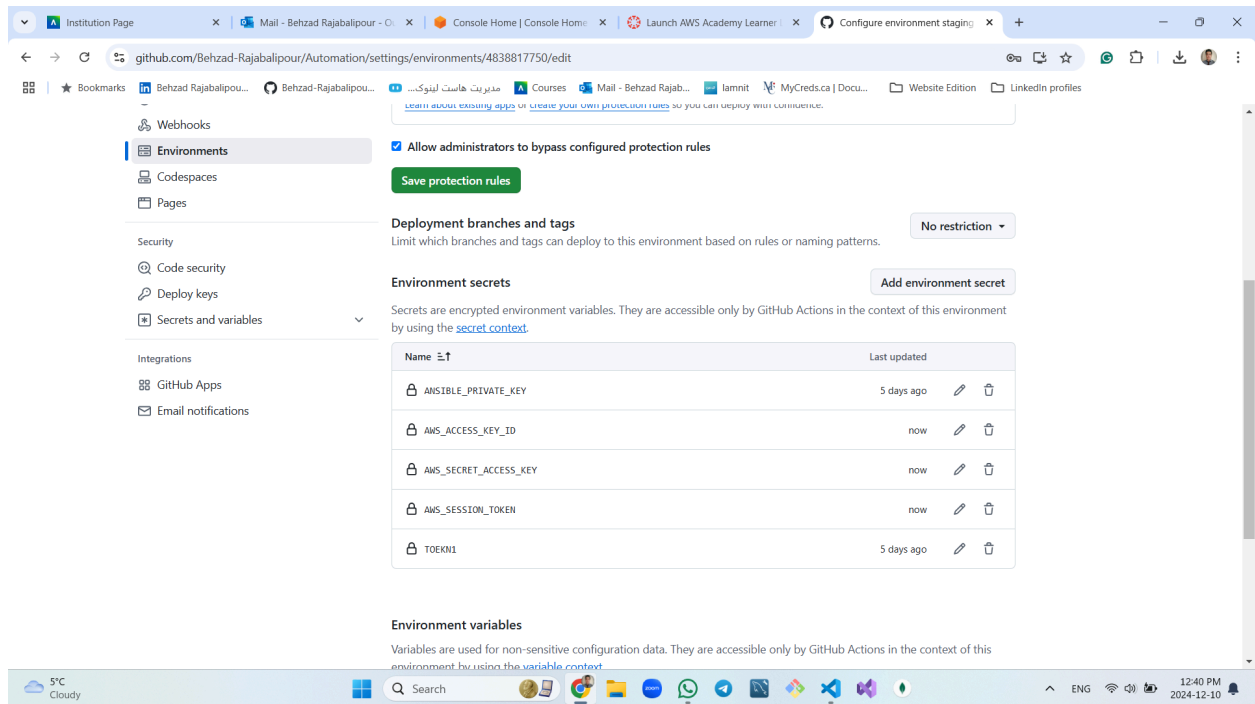
Generate a webpage using jinja2 and other configurations to the webpage

## GitHub actions

```

jobs:
  terraform:
    name: "Terraform - Network and Webserver"
    runs-on: ubuntu-latest
    environment: staging
    env:
      GITHUB_TOKEN: ${ secrets.GITHUB_TOKEN }
      AWS_ACCESS_KEY_ID: ${ secrets.AWS_ACCESS_KEY_ID }
      AWS_SECRET_ACCESS_KEY: ${ secrets.AWS_SECRET_ACCESS_KEY }
      AWS_SESSION_TOKEN: ${ secrets.AWS_SESSION_TOKEN }

```



## Providing the access credentials to GitHub

```
# --- Network Infrastructure ---
- name: Terraform Format (Network)
  id: fmt-network
  run: terraform fmt
  working-directory: ./Terraform/prod/Network
  continue-on-error: true

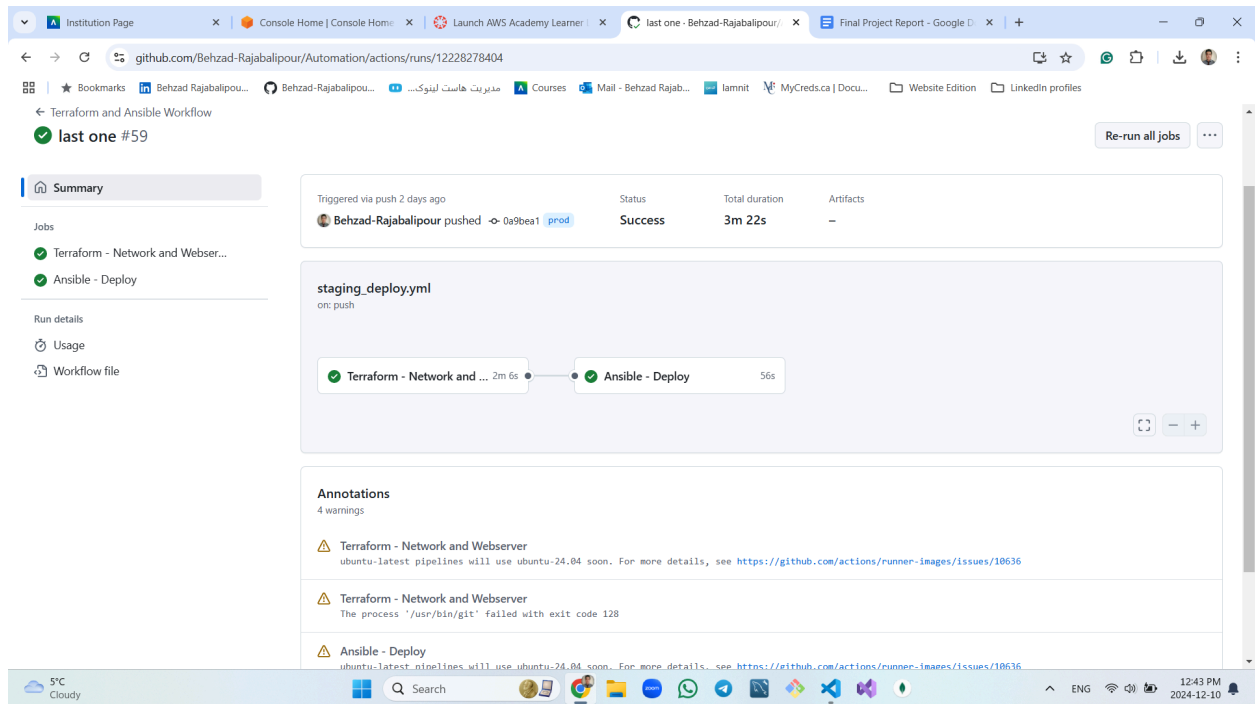
- name: Terraform Init (Network)
  id: init-network
  run: terraform init
  working-directory: ./Terraform/prod/Network

- name: Terraform Validate (Network)
  id: validate-network
  run: terraform validate -no-color
  working-directory: ./Terraform/prod/Network

- name: Terraform Plan (Network)
  id: plan-network
  run: terraform plan -input=false -no-color -out=tf.plan
  working-directory: ./Terraform/prod/Network

- name: Terraform Apply (Network)
  id: apply-network
  run: terraform apply -input=false tf.plan
  working-directory: ./Terraform/prod/Network
```

sequence of steps to format, initialize, validate, plan, and apply Terraform configurations for the network infrastructure



Picture shows Both Terraform and Ansible part installed correctly

```
# --- Webserver Infrastructure ---
- name: Terraform Format (Webserver)
  id: fmt-webserver
  run: terraform fmt
  working-directory: ./Terraform/prod/Webserver
  continue-on-error: true

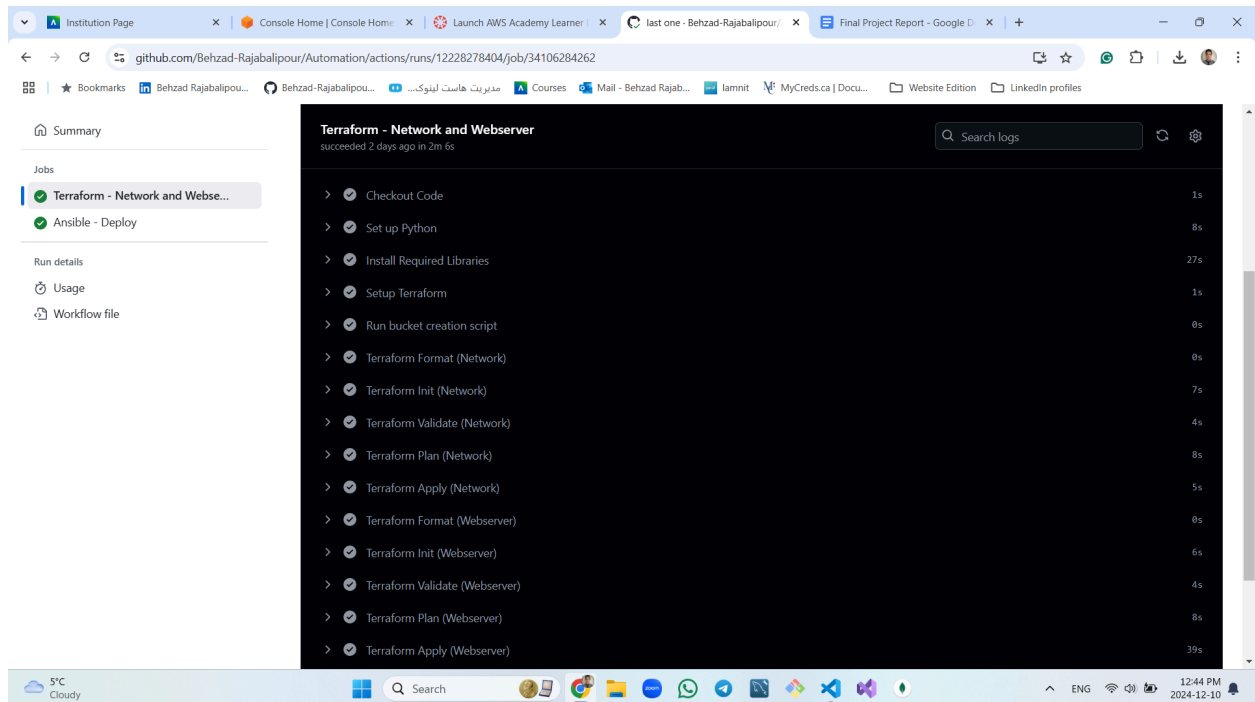
- name: Terraform Init (Webserver)
  id: init-webserver
  run: terraform init
  working-directory: ./Terraform/prod/Webserver

- name: Terraform Validate (Webserver)
  id: validate-webserver
  run: terraform validate -no-color
  working-directory: ./Terraform/prod/Webserver

- name: Terraform Plan (Webserver)
  id: plan-webserver
  run: terraform plan -input=false -no-color -out=tf.plan
  working-directory: ./Terraform/prod/Webserver

- name: Terraform Apply (Webserver)
  id: apply-webserver
  run: terraform apply -input=false tf.plan
  working-directory: ./Terraform/prod/Webserver
```

sequence of steps to format, initialize, validate, plan, and apply Terraform configurations for the webserver infrastructure



Github Action workflow deployed both Network and Webserver directory

```
ansible:
  name: "Ansible - Deploy"
  runs-on: ubuntu-latest
  needs: terraform
  environment: staging

  env:
    GITHUB_TOKEN: ${ secrets.GITHUB_TOKEN }
    AWS_ACCESS_KEY_ID: ${ secrets.AWS_ACCESS_KEY_ID }
    AWS_SECRET_ACCESS_KEY: ${ secrets.AWS_SECRET_ACCESS_KEY }
    AWS_SESSION_TOKEN: ${ secrets.AWS_SESSION_TOKEN }

  steps:
    - name: Checkout Code
      uses: actions/checkout@v3

    - name: Set up Python (for Ansible)
      uses: actions/setup-python@v2
      with:
        python-version: "3.10.8"

    - name: Install Required Libraries
      run: |
        python -m pip install --upgrade pip
        pip install ansible==2.9 boto3 botocore "Jinja2<3.1"

    - name: Verify Inventory File
      run: |
        cat ./ansible/inventories/aws_ec2.yaml

    - name: Add SSH Key
      # it will store in the GitHub Action Ubuntu server
```

GitHub Actions workflow to deploy infrastructure using Ansible

The steps here include setting up Python and installing required libraries verifying the Ansible inventory and SSH key which will be used to access the targeted EC2 instances for configuration

```
Source: github-workflow.json
- name: Run Ansible Playbook
  run: |
    ansible-playbook -i ./ansible/inventories/aws_ec2.yaml ./ansible/playbook_jinja2.yaml
  env:
    ANSIBLE_HOST_KEY_CHECKING: "false"
```

Run the Ansible playbook

## Deployment

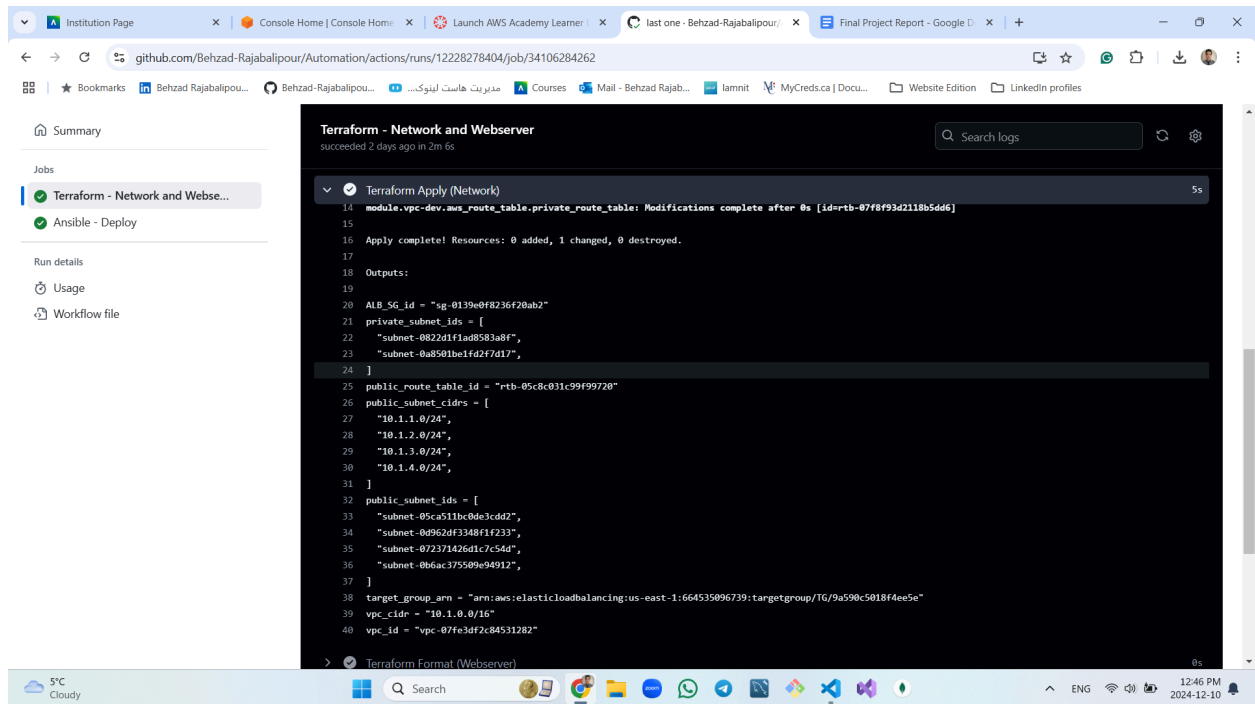
Here are some screenshots of the successful deployment of the website

```
Apply complete! Resources: 0 added, 1 changed, 0 destroyed.


Outputs:

ALB_SG_id = "sg-007153cab24d1ede1"
private_subnet_ids = [
  "subnet-077a407460760a471",
  "subnet-06f990e791c2db207",
]
public_route_table_id = "rtb-0278ff3538f0c061f"
public_subnet_cidrs = [
  "10.1.1.0/24",
  "10.1.2.0/24",
  "10.1.3.0/24",
  "10.1.4.0/24",
]
public_subnet_ids = [
  "subnet-0d5e170f2aa801f89",
  "subnet-07300ba78ba94fb77",
  "subnet-08ad85a07b606657f",
  "subnet-09e45886882ae142f",
]
target_group_arn = "arn:aws:elasticloadbalancing:us-east-1:142390386045:targetgroup/TG/e42fed72659dbce7"
vpc_cidr = "10.1.0.0/16"
vpc_id = "vpc-0ee5ec7cc67280da5"
voclabs:~/environment/project/Terraform/prod/Network $
```

Networking deployment in Terraform

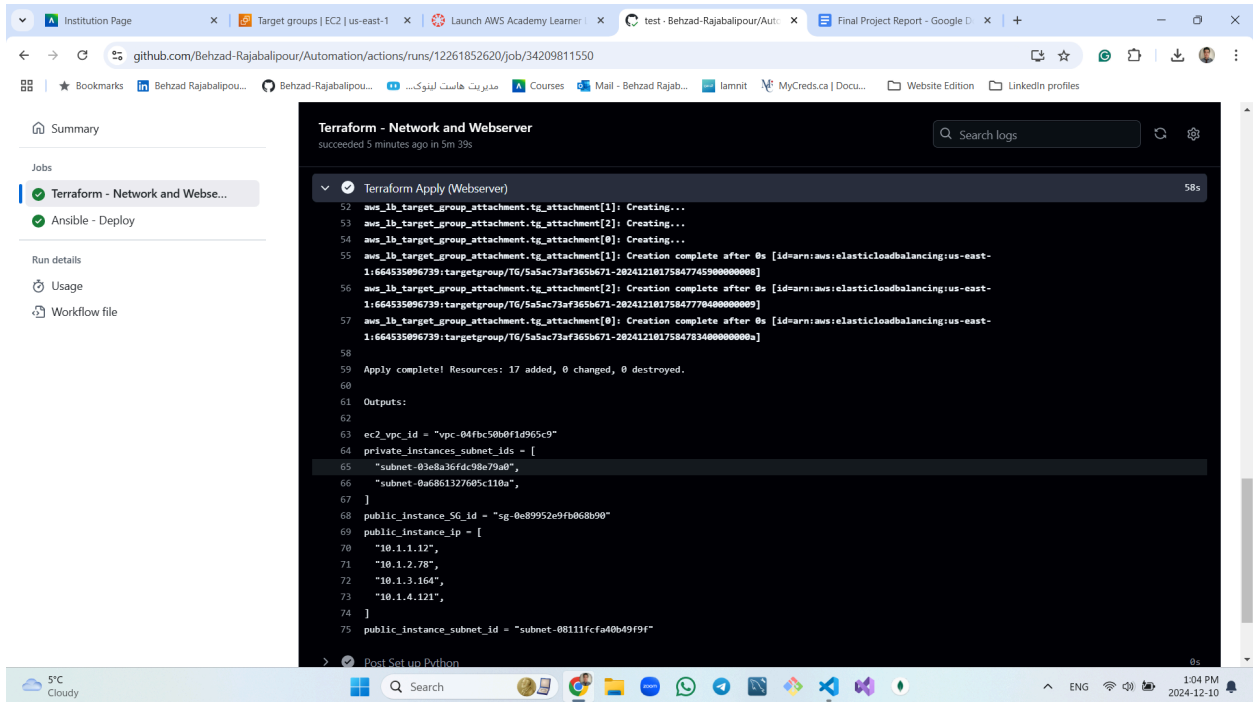


## Networking deployment in GitHub Action




## Web server deployment(EC2 and installing httpd) in Terraform






## Web server deployment(EC2 and installing httpd) in Anisble




this is Ammar and Behzad, My private IP is 10.1.3.137 in PROD environment

Built by Terraform!



this is Ammar and Behzad, My private IP is 10.1.1.186 in PROD environment

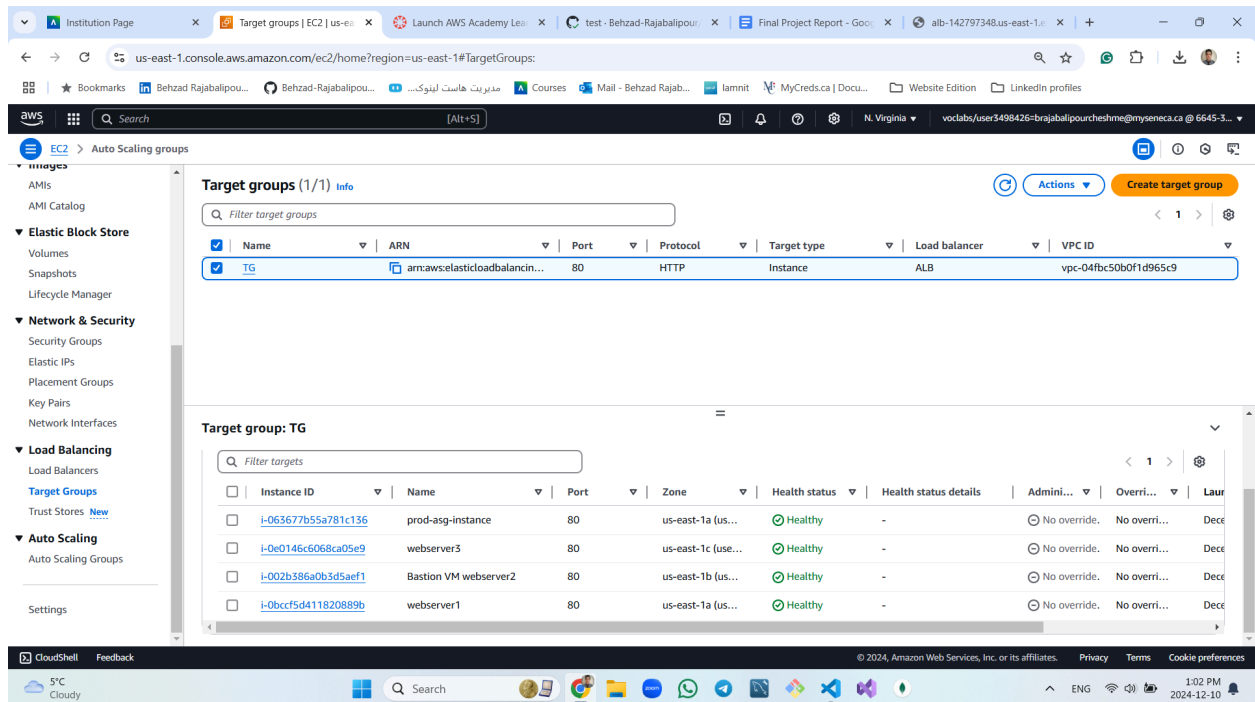
Built by Terraform!



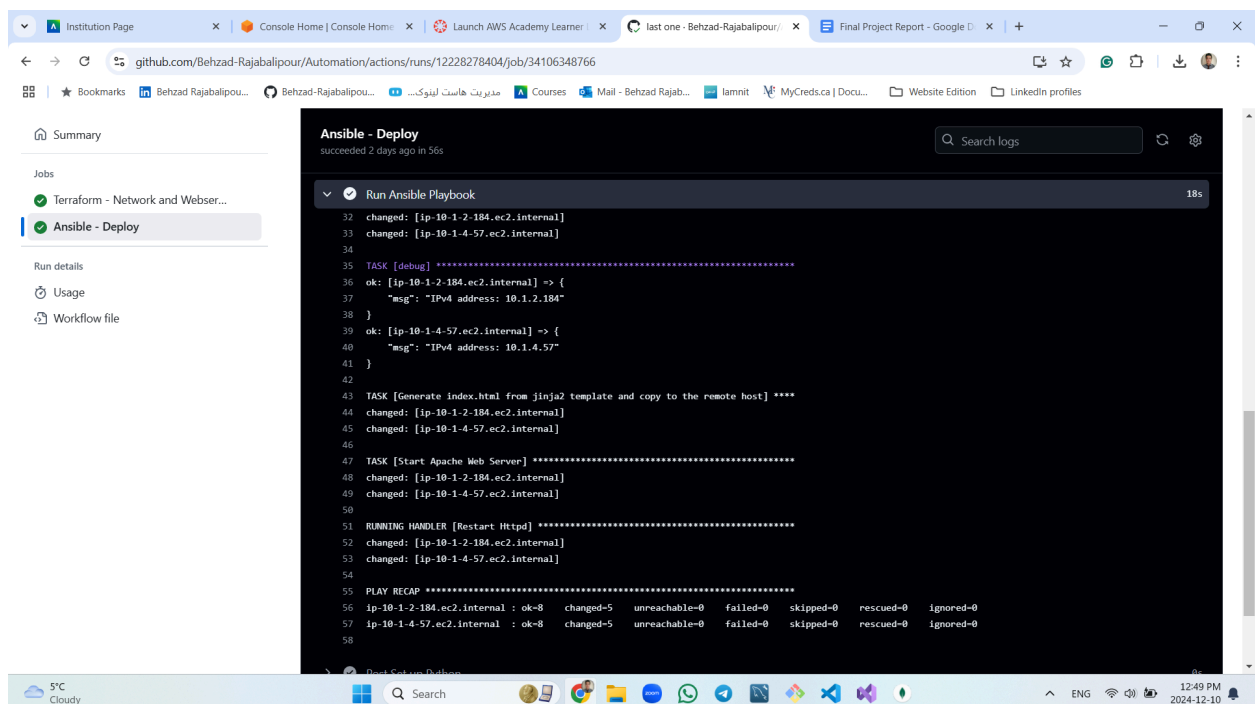
"Use new template in 10.1.2.78"

This server runs ip-10-1-2-78,Amazon 2

This is Webserver2 (Bastion Host) which installed http with Ansible



The target group at the final step, shows all 3 instances are healthy. One more instance (prod-asg-instance) which is created by ASG  
Accessing multiple instances via the Load balancer



Webserver deployment in GitHub workflow (Run Ansible)

Ansible - Deploy  
succeeded 2 days ago in 56s

Jobs

- Terraform - Network and Webser...
- Ansible - Deploy

Run details

Usage

Workflow file

```
1 ▶ Run cat ./ansible/inventories/aws_ec2.yaml
2 # ansible-inventory -i inventories/aws_ec2.yaml --graph
3 ---
4 plugin: aws_ec2
5 regions:
6   - us-east-1
7 keyed_groups:
8   - key: tags.Type      # Groups instances by their Type tag, like Type = linux or Type = windows
9     prefix: tag         # Adds the prefix tag_ to the group names, like tag_linux
10 filters:
11   instance-state-name : running
12 compose:
13   ansible_host: public_ip_address  # ansible_host: is mandatory like ansible_user: in tag_type1.yaml
14   ansible_hostname: public_ip_address  # i used it in index.j2 file
15   # ansible_user: ec2-user
16   # ansible_ssh_private_key_file: ~/.ssh/prodKey  # it will store in the Github Action Ubuntu server
17
18 ▶ Add SSH Key
19
20 ▶ Run Ansible Playbook
```

## Ansible-Inventory file

us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#Instances:

EC2 > Auto Scaling groups

Dashboard

EC2 Global View

Events

Instances

- Instances
- Instance Types
- Launch Templates
- Spot Requests
- Savings Plans
- Reserved Instances
- Dedicated Hosts
- Capacity Reservations

Images

- AMIs
- AMI Catalog

Elastic Block Store

- Volumes
- Snapshots
- Lifecycle Manager

Network & Security

- Security Groups
- Elastic IPs

Instances (1/7) Info

Find Instance by attribute or tag (case-sensitive)

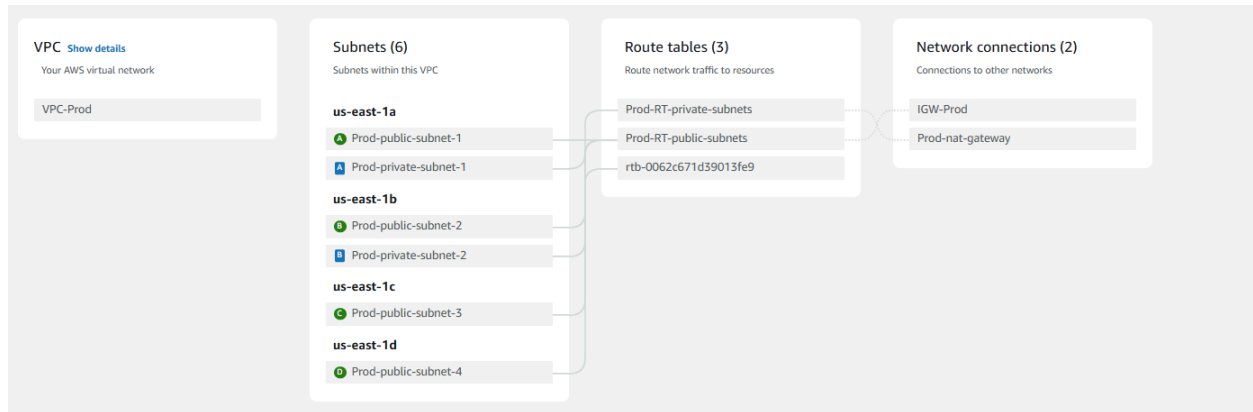
Running

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS
Bastion VM webserver2	i-002b386a0b3d5aef1	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	-
VM6	i-04928281058538f6a	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1b	-
webserver3	i-0e0146c6068ca05e9	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1c	-
prod-asg-instance	i-063677b55a781c136	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1a	-
webserver1	i-0bccf5d411820889b	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1a	-
Webserver5	i-0caa23216a2b0a114	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1a	-
webserver4	i-01b36bec4f08fbb4	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1d	-

i-01b36bec4f08fbb4 (webserver4)

Key	Value
env	prod
Owner	Behzad
Name	webserver4
Type	type1
App	FinalProject

4 instances in public subnets, 2 instances in private subnets



VPC with 4 availability zones

## Challenges

The main challenge was coordination between the team members. As each step is dependent on the previous step, we had to wait and then build the next based on what we created. GitHub should have streamlined this process but with the limited time we had, it added more overhead. We eventually decided to divide up the project into three main steps: Terraform, Ansible, and GitHub actions. The real challenge came in integrating everything which felt like solving a puzzle.

One more challenge we got was about ASG. The Auto Scaling Group would create a new instance without considering how many instances are already inside the Target Group. We even used the same templates for webserver but ASG still create a new instance without considering there were enough instances inside the Target group.

## Conclusion

In conclusion, this project demonstrates the seamless integration of Terraform, Ansible, and GitHub Actions to automate the deployment and configuration of a highly available website.

Terraform for infrastructure provisioning, Ansible for server configuration, and GitHub Actions for continuous integration